# Reproducing “Near-Duplicate Detection for Web Crawling”

ToDo

2026-08-01

**Abstract** This paper aims to reproduce the findings of the “Near-Duplicate Detection for Web Crawling” work, using PySpark to scale the proposed algorithm, under different datasets, feature models and K values.

## Introduction

Two documents are considered near-duplicates if they are similar to each other. An efficient and optimized algorithm for near-duplicates detection in a dataset is crucial for tasks like web crawling and, recently, cleaning datasets for AI models training. Manku et al. paper suggests to extract features from documents, compute the aggregated fingerprint using Charikar’s Simhash and lastly compare fingerprints looking only at a specific partition of bits, for every partition. This paper reproduces the central claims of [1] and analyzes its performance under different datasets, feature models and K values.

## Related work

This is the related work section.

## Setup

This is the setup section.

## Methodology

This is the methodology section.

## Results

This is the results section.

## Discussion

This is the discussion section.

## Conclusion

This is the conclusion section.

## Bibliography

[1] A. D. S. Gurmeet Singh Manku Arvind Jain, “Detecting Near-Duplicates for Web Crawling,” 2007.